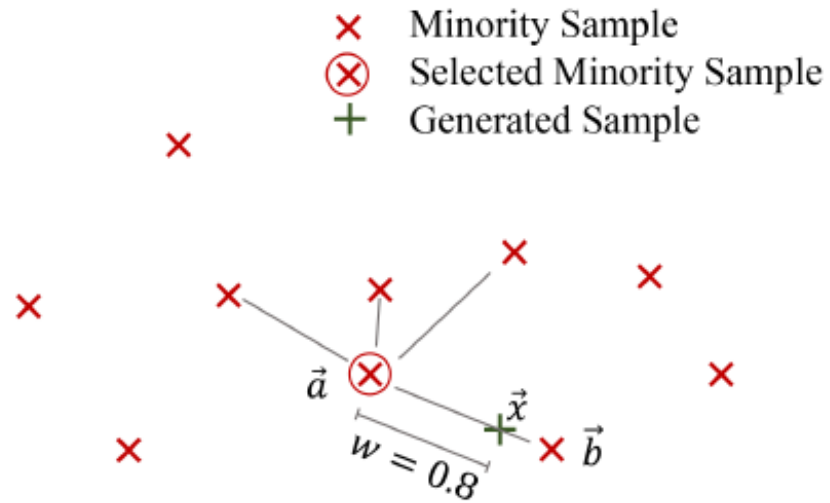




K-Means SMOTE

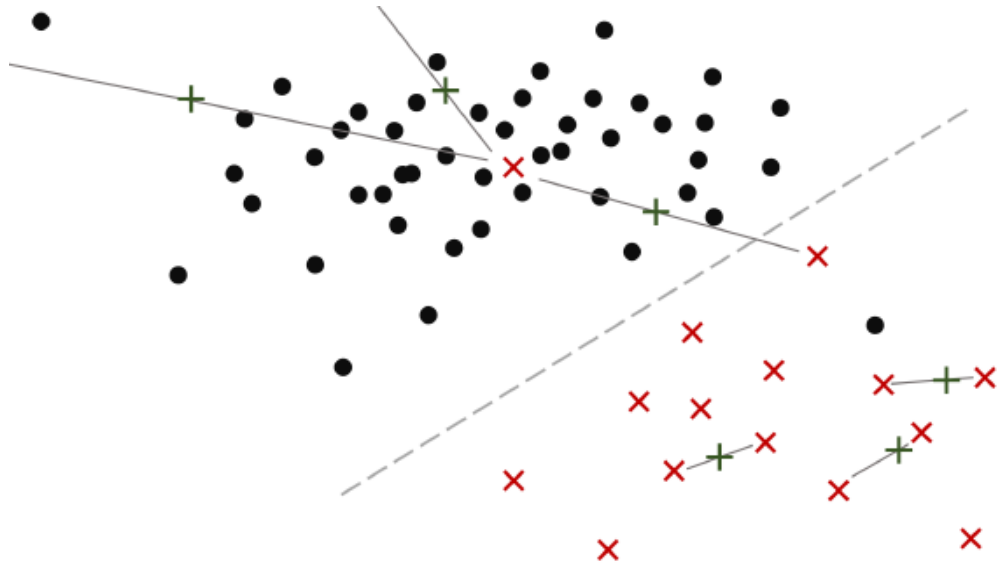
A bit of background: SMOTE



<https://arxiv.org/pdf/1711.00837.pdf>

SMOTE linearly interpolates a randomly selected minority sample and one of its $k = 5$ nearest neighbors


SVM, Borderline SMOTE



<https://arxiv.org/pdf/1711.00837.pdf>

- We should not create samples in areas that are safe
- We should create samples at the boundary

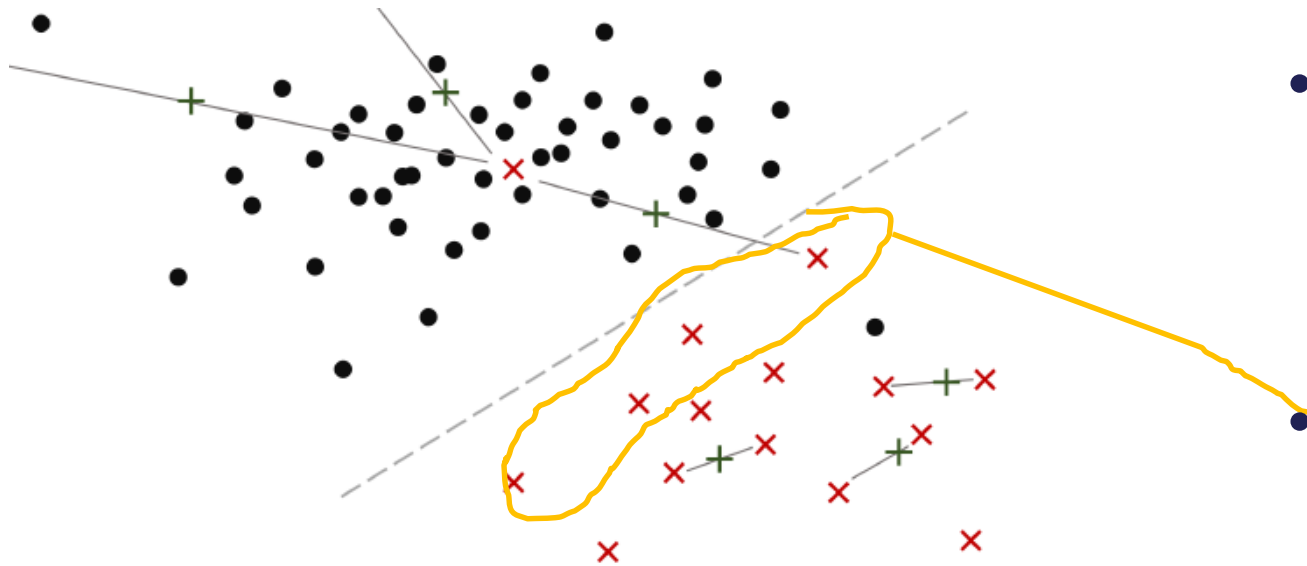
SVM, Borderline SMOTE



<https://arxiv.org/pdf/1711.00837.pdf>

- We should not create samples in areas that are safe
- We should create samples at the boundary

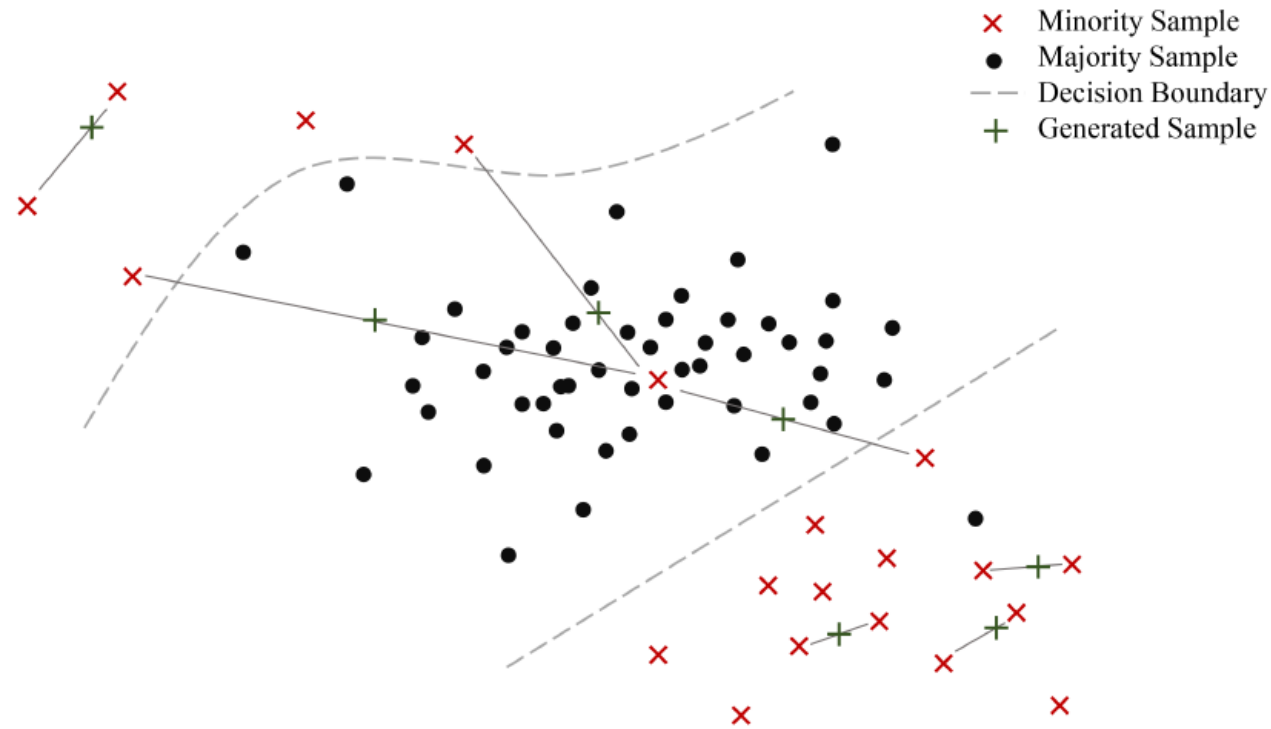
SVM, Borderline SMOTE



<https://arxiv.org/pdf/1711.00837.pdf>

- We should not create samples in areas that are safe
- We should create samples at the boundary

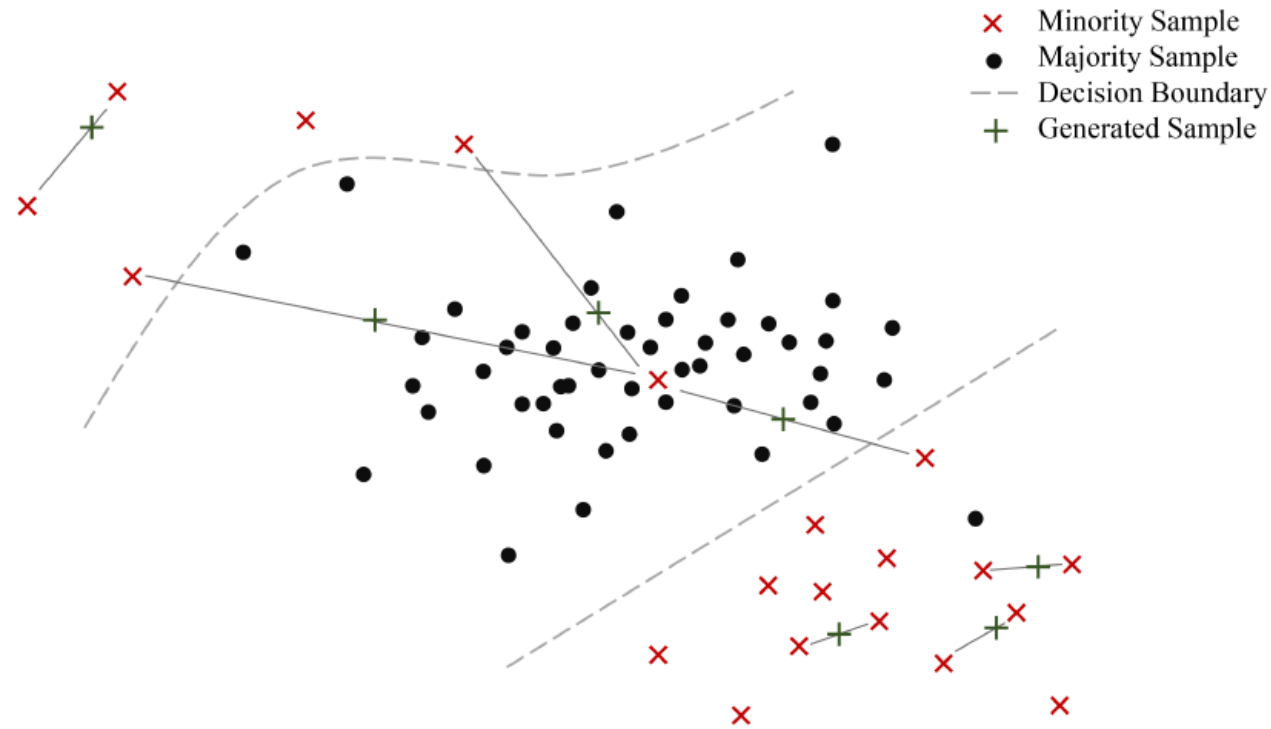
SMOTE, noise and intra-class clusters



<https://arxiv.org/pdf/1711.00837.pdf>

- SMOTE might create noisy examples
- SMOTE does not contemplate intra-class clusters

Examples of intra-class clusters



<https://arxiv.org/pdf/1711.00837.pdf>

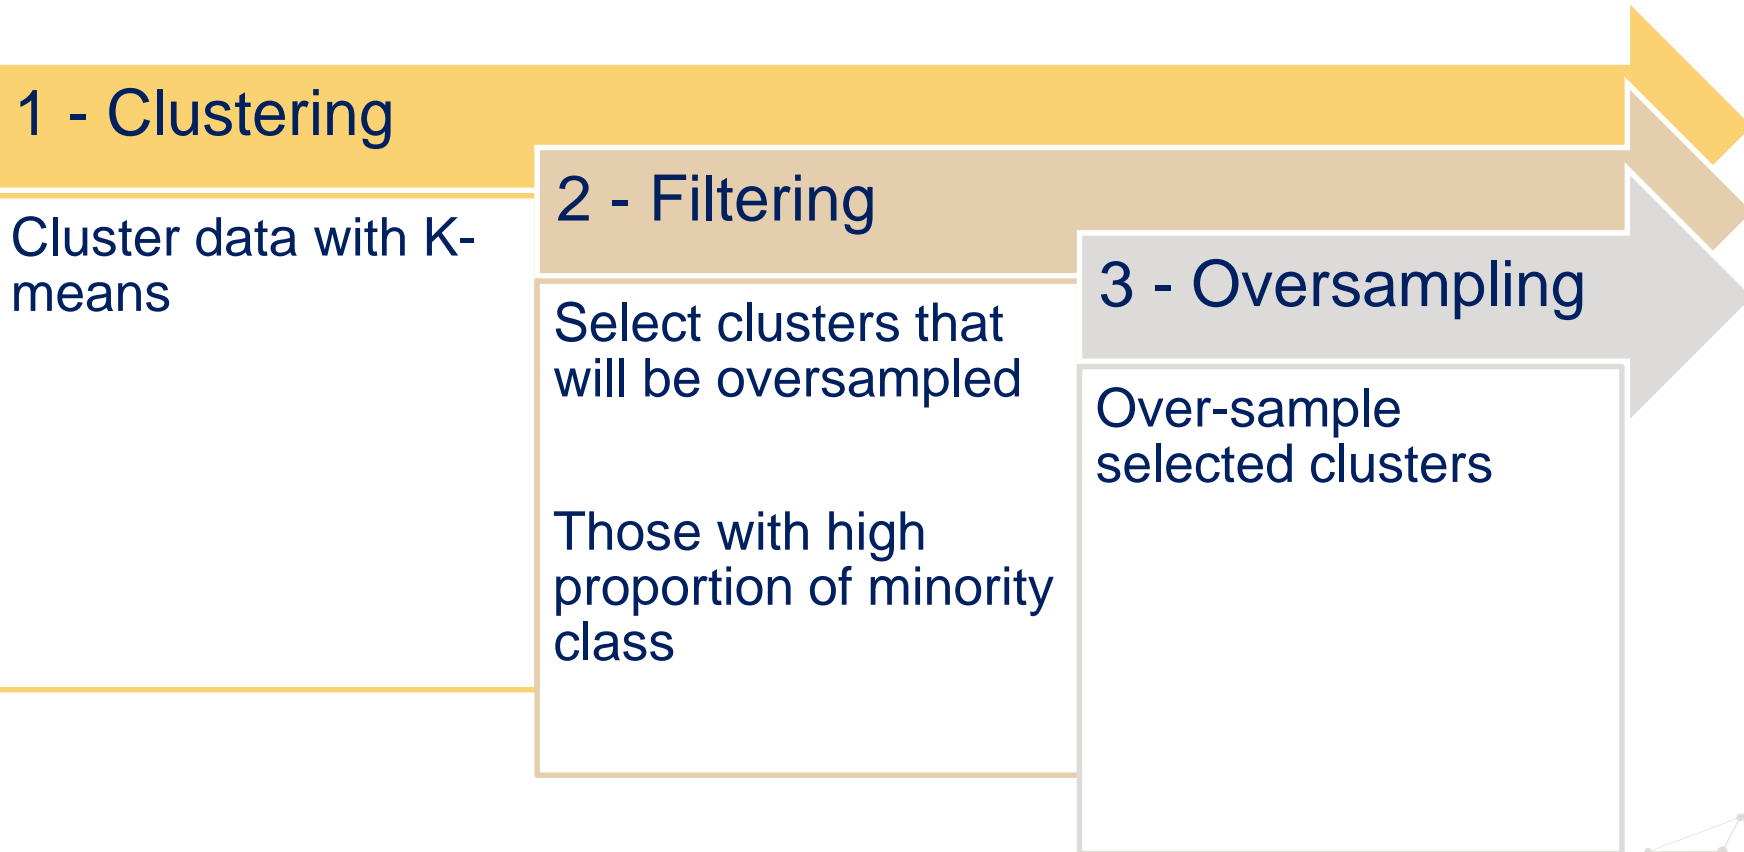
Fraudulent credit card transactions:

- Transactions in foreign countries
- High value transactions


K-Means SMOTE – the idea

- Boost minority class regions by creating samples within naturally occurring clusters of the minority class.
- Contemplate intra-class clusters
- Avoid introducing noise


K-Means SMOTE - 3 steps



Input data




Find $k = 3$ clusters and compute imbalance ratio (IR)




- ✕ Minority Sample
- Majority Sample
- - Decision Boundary
- + Generated Sample

Oversampled data rectifies decision boundary



Use SMOTE to oversample clusters with $IR > 1$, generating more samples in sparse clusters




<https://arxiv.org/pdf/1711.00837.pdf>

Diagram showing
the 3 steps of K-
means SMOTE

By default
clusters where
50% are minority
are selected by
default

Clustering step




Find clusters → K means algorithm with entire dataset

We need to know K, or treat K as a hyperparameter

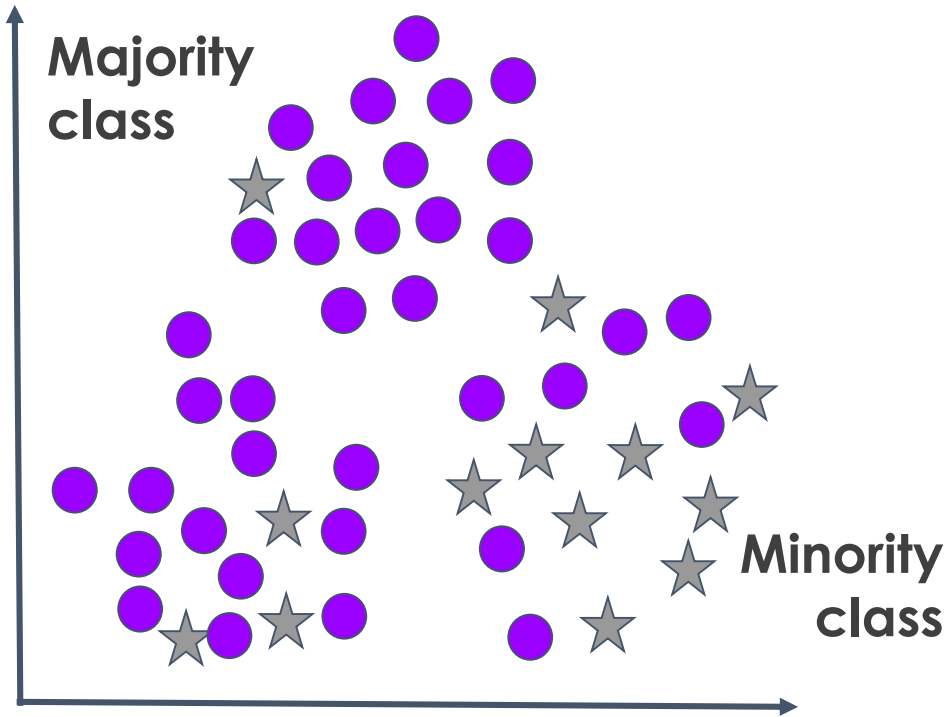
K-means may be slow to converge in huge datasets

Filtering step




- By default select those clusters where 50% of observations belong to minority
- $IR = 1 = \# \text{ minority} / \# \text{ majority}$
- We can increase IR, thus we select clusters with higher proportion of minority class
- IR becomes another hyperparameter

Over-sampling



- Determine how many samples to create in each cluster
- Assign weights to clusters, more weights to clusters with less minority observations.

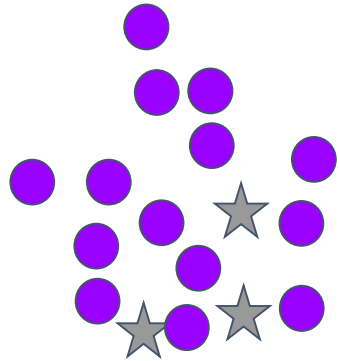
Over-sampling



We take this cluster to proceed with the demo

- Determine how many samples to create in each cluster
- Assign weights to clusters, more weights to clusters with less minority observations.

Cluster weight



1. Determine Euclidean distance between all samples from minority
2. Determine Mean Euclidean distance
3. $\text{density} = \frac{X_{min}}{L2_{mean} \text{ number of features}}$
3. Sparsity = 1 / density
4. Cluster sparsity = Sparsity / sum(Sparsity all clusters)

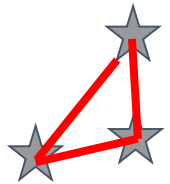
Cluster weight

1. Determine Euclidean distance between all samples from minority
2. Determine Mean Euclidean distance
3. $\text{density} = \frac{X_{\min}}{L2_{\text{mean}}^{\text{number of features}}}$
3. $\text{Sparsity} = 1 / \text{density}$
4. $\text{Cluster sparsity} = \text{Sparsity} / \text{sum}(\text{Sparsity all clusters})$



Cluster weight

1. Determine Euclidean distance between all samples from minority
2. Determine Mean Euclidean distance (L2-mean)
3.
$$\text{density} = \frac{\# \text{ minority}}{L2\text{mean} \text{ number of features}}$$
3. Sparsity = 1 / density
4. Cluster sparsity = Sparsity / sum(Sparsity all clusters)



K-Means SMOTE

- Calculate the number of synthetic examples that need to be generated for each cluster


$$g_i = csi \times G$$

csi = cluster sparsity

G = total number of samples to generate

g_i = number of samples to generate from cluster i

SMOTE with samples in cluster



<https://arxiv.org/pdf/1711.00837.pdf>

- Linearly interpolate between sample and neighbour chosen at random
- If cluster has few samples, we may not have enough neighbours
- The number of neighbours becomes a hyperparameter

K-mean SMOTE – my thoughts

- The rationale is well thought, it makes sense
- The implementation of the algorithm is not super straight forward
- A lot of parameters to adjust
- Potentially some EDA to corroborate those parameters

Imbalanced-learn: KMeansSMOTE

```
▶ sm = KMeansSMOTE(  
    sampling_strategy='auto',  # samples only the minority class  
    random_state=0,  # for reproducibility  
    k_neighbors=2,  
    n_jobs=None,  
    kmeans_estimator=KMeans(n_clusters=3, random_state=0),  
    cluster_balance_threshold=0.1,  
    density_exponent='auto'  
)  
  
X_res, y_res = sm.fit_resample(X, y)
```

THANK YOU

www.trainindata.com